

Lab 3A: Implement Simple Linear Regression

Import Libraries

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Import Dataset

In [2]:

```
dataset = pd.read_csv("Salary.csv")
```

In [3]:

```
dataset.head()
```

Out[3]:

| | YearsExperience | Salary |
|---|-----------------|---------|
| 0 | 1.1 | 39343.0 |
| 1 | 1.3 | 46205.0 |
| 2 | 1.5 | 37731.0 |
| 3 | 2.0 | 43525.0 |
| 4 | 2.2 | 39891.0 |

In [4]:

```
dataset.shape
```

Out[4]:

```
(30, 2)
```

In [5]:

```
dataset.columns
```

Out[5]:

```
Index(['YearsExperience', 'Salary'], dtype='object')
```

In [6]:

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   YearsExperience  30 non-null     float64
 1   Salary          30 non-null     float64
dtypes: float64(2)
memory usage: 608.0 bytes
```

Preprocessing steps

In [7]:

```
# Step 1: Divide dataframe into i/p independent variable and output dependent features
X= dataset.iloc[:, :-1]
Y= dataset.iloc[:, -1]
```

In [8]:

```
print(X.head())
```

```
   YearsExperience
0              1.1
1              1.3
2              1.5
3              2.0
4              2.2
```

In [9]:

```
print(Y.head())
```

```
0    39343.0
1    46205.0
2    37731.0
3    43525.0
4    39891.0
Name: Salary, dtype: float64
```

In [10]:

```
# Step 4: Splitting Data into Training and Testing
from sklearn.model_selection import train_test_split

X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=3)
```

In [11]:

```
print(X_train.shape)
print(X_test.shape)
```

(24, 1)

(6, 1)

Create Regression model on training data

In [12]:

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train,Y_train)
```

Out[12]:

LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [13]:

```
print("Intercept: ", regressor.intercept_)
print("Coefficient: ", regressor.coef_)
```

Intercept: 26197.76702738423

Coefficient: [9418.56697369]

From above values our model regression line equation

$y=b_0+b_1*X$

$y=26197.767027384238 + 9418.56697369*X$

by using above equation lwets find salary for experience=6

In [14]:

```
ynew=26197.767027384238 + 9418.56697369*6
print("Salary for 6 yr exp: ",ynew)
```

Salary for 6 yr exp: 82709.16886952423

In [15]:

```
ypred =regressor.predict(X_test)
```

In [16]:

```
print(ypred)
print(Y_test)
```

```
[ 72348.74519846  53511.61125108 100604.44611953 115674.15327744
 81767.31217215  68581.31840899]
15      67938.0
5       56642.0
22     101302.0
26     116969.0
18      81363.0
14      61111.0
Name: Salary, dtype: float64
```

Accuracy Of Model

In [17]:

```
# 1. training accuracy

print("Training Accuracy: ",regressor.score(X_train,Y_train))
```

Training Accuracy: 0.954524919039405

In [18]:

```
# 2. testing accuracy

print("Testing Accuracy: ",regressor.score(X_test,Y_test))
```

Testing Accuracy: 0.9695039421049821

Visualizing the model

Visualize the training data

In [19]:

```
plt.scatter(X_train, Y_train, color = "red" ,s=12)
plt.plot(X_train, regressor.predict(X_train), color = "blue")
plt.title("Simple Regressor on Training Data")
plt.xlabel("Experience in years")
plt.ylabel("Salary")
plt.show()
```



In [20]:

```
plt.scatter(X_test, Y_test, color = "red" ,s=12)
plt.plot(X_test, regressor.predict(X_test), color = "blue")
plt.title("Simple Regressor on Testing Data")
plt.xlabel("Experience in years")
plt.ylabel("Salary")
plt.show()
```



Find R² Score of model

In [21]:

```
from sklearn.metrics import r2_score  
r2_score(Y_test,ypred)
```

Out[21]:

0.9695039421049821

In []:

Test Your Knowledge

Q] Predict Salary for 6.1 Year Experience

In [29]:

```
pred = regressor.predict([[6.1]])  
pred
```

C:\Users\Admin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(

Out[29]:

array([83651.02556689])

In [31]:

```
predi=26197.767027384238 + 9418.56697369*6.1  
predi
```

Out[31]:

83651.02556689324

Q] Change the random state

In [32]:

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=57)

# Linear Regression model
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train,Y_train)
```

Out[32]:

LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [33]:

```
# model on testing data
ypred = regressor.predict(X_test)
ypred
```

Out[33]:

```
array([82225.47392205, 45871.84304206, 65446.87505436, 63582.58629129,
       57057.57562052, 69175.45258052])
```

In [34]:

```
# Accuracy, R^2
print("Training accuracy: ",regressor.score(X_train,Y_train) )
print("Testing accuracy: ",regressor.score(X_test,Y_test) )
print("R^2 score: ",r2_score(Y_test,ypred) )
```

```
Training accuracy:  0.9576054599427444
Testing accuracy:  0.809635146935763
R^2 score:  0.809635146935763
```

In [35]:

```
# Predict Salary for 6.1 Year Experience
regressor.predict([[6.1]])
```

```
C:\Users\Admin\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

Out[35]:

```
array([84089.76268513])
```

In []:

