

Lab 3B: Implement multiple linear regression

Import Libraries

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Import Dataset

In [2]:

```
dataset = pd.read_csv("50_Startups.csv")
```

In [3]:

```
dataset.head()
```

Out[3]:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

In [4]:

```
dataset.shape
```

Out[4]:

(50, 5)

In [5]:

```
dataset.columns
```

Out[5]:

Index(['R&D Spend', 'Administration', 'Marketing Spend', 'State', 'Profit'], dtype='object')

In [6]:

dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   R&D Spend              50 non-null    float64
1   Administration         50 non-null    float64
2   Marketing Spend        50 non-null    float64
3   State                  50 non-null    object
4   Profit                 50 non-null    float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

Step 1: Divide dataframe into independent variable/ input and dependent / output features

In [7]:

```
X = dataset.iloc[:, :-1]
Y = dataset.iloc[:, -1]
```

In [8]:

```
print(X.head())
print(Y.head())
```

	R&D Spend	Administration	Marketing Spend	State
0	165349.20	136897.80	471784.10	New York
1	162597.70	151377.59	443898.53	California
2	153441.51	101145.55	407934.54	Florida
3	144372.41	118671.85	383199.62	New York
4	142107.34	91391.77	366168.42	Florida

```
0    192261.83
1    191792.06
2    191050.39
3    182901.99
4    166187.94
```

Name: Profit, dtype: float64

Encoding the Categorical Data

In [9]:

dataset['State'].value_counts()

Out[9]:

```
New York      17
California    17
Florida       16
Name: State, dtype: int64
```

In [10]:

```
# as we have 3 categories
# so we use one hot encoder
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

ct= ColumnTransformer(transformers = [('encoder', OneHotEncoder(),[3])],remainder='passt
X=np.array(ct.fit_transform(X))
```

In [28]:

```
# print(X)
```

Step 4: Split the data into training and testing

In [12]:

```
from sklearn.model_selection import train_test_split
X_train, X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=1)
```

In [13]:

```
print(X_train.shape)
print(X_test.shape)
```

```
(40, 6)
(10, 6)
```

Lets test linear regression model on training data

In [14]:

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train,Y_train)
```

Out[14]:

```
LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [15]:

```
print("coefficient for state=California (B1) = ", regressor.coef_[0])
print("coefficient for state=Florida (B2) = " ,regressor.coef_[1])
print("coefficient for state=New York (B3) = ", regressor.coef_[2])
print("coefficient for R&D (B4) = " ,regressor.coef_[3])
print("coefficient for Adminstration (B5) = " ,regressor.coef_[4])
print("coefficient for Marketing Spent (B6) = ", regressor.coef_[5])
print("Initial Profit Intercept (B0) = " ,regressor.intercept_)
```

```
coefficient for state=California (B1) = -285.1777694629406
coefficient for state=Florida (B2) = 297.56087646106005
coefficient for state=New York (B3) = -12.383107002767183
coefficient for R&D (B4) = 0.7743420811125808
coefficient for Adminstration (B5) = -0.009443695851296884
coefficient for Marketing Spent (B6) = 0.028918313285055255
Initial Profit Intercept (B0) = 49834.885073215744
```

Predicting the test result

In [16]:

```
ypred = regressor.predict(X_test)
```

In [17]:

```
print(ypred)
print(Y_test)
```

```
[114664.41715867  90593.1553162   75692.84151574  70221.88679651
 179790.25514874 171576.92018522  49753.58752029 102276.65888936
 58649.37795761  98272.02561131]
27    105008.31
35     96479.51
40     78239.91
38     81229.06
2     191050.39
3     182901.99
48     35673.41
29     101004.64
46     49490.75
31     97483.56
Name: Profit, dtype: float64
```

Testing the single instance

For example the profit of a startup with R&D spend=160000, Adminstration Spend= 130000, Marketing Spend = 300000 State='California'

I/P : [1,0,0, 160000, 130000, 300000]

In [18]:

```
regressor.predict([[1,0,0, 160000, 130000, 300000]])  
# [1,0,0,160.....] == 1,0,0 = California, Florida, New York
```

Out[18]:

```
array([180892.25380661])
```

Accuracy of the model

In [19]:

```
# 1. Training Accuracy  
tr = regressor.score(X_train,Y_train)  
print("Training Accuracy = ", tr)
```

```
Training Accuracy = 0.942446542689397
```

In [20]:

```
# 1. Test Accuracy  
te = regressor.score(X_test,Y_test)  
print("Testing Accuracy = ", te)
```

```
Testing Accuracy = 0.9649618042060673
```


Build Residual Error Plot

In [21]:

```
# Setting plot style
plt.style.use('fivethirtyeight')

plt.scatter(regressor.predict(X_train), regressor.predict(X_train)-Y_train, color = 'green')
plt.hlines(y=0, xmin=0,xmax=200000,linewidth=2)

# Plotting residual errors in testing data
plt.scatter(regressor.predict(X_test), regressor.predict(X_test)-Y_test, color = 'red',
plt.hlines(y=0, xmin=0,xmax=200000,linewidth=2)
plt.legend(['Training Data', 'Testing Data'],loc='upper right')
plt.title("Residual Error")
plt.xlabel("Profit")
plt.ylabel("Error")
plt.show()
```



R² Score

In [22]:

```
from sklearn.metrics import r2_score
r2_score(Y_test,ypred)
```

Out[22]:

0.9649618042060673

In [23]:

```
r2_score(regressor.predict(X),Y)
```

Out[23]:

0.9439977800352345

In [25]:

```
print("Intercept: ", regressor.intercept_)  
print("Coefficient: ", regressor.coef_)
```

```
Intercept: 49834.885073215744  
Coefficient: [-2.85177769e+02  2.97560876e+02 -1.23831070e+01  7.74342081  
e-01  
-9.44369585e-03  2.89183133e-02]
```

In []:

Test Your Knowledge

Q] Predict profit for { R&D Spend = 300000, Administration Spend = 160000, Marketing Spend = 200000, State = Florida }

In [26]:

```
regressor.predict([[0,1,0,300000,160000,200000]])
```

Out[26]:

```
array([286707.74160425])
```

Q] Change the random state

In [33]:

```
# splitting into training and testing  
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=57)  
  
# linear regression model  
from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(X_train, Y_train)
```

Out[33]:

```
LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [38]:

```
# model on testing data
ypred = regressor.predict(X_test)
ypred
```

Out[38]:

```
array([101665.54137004, 115967.55312425, 116484.53520105, 154483.34482823,
       82520.71288925,  67442.66965265,  99153.93967286,  98982.68177633,
       134913.84345446,  96848.151744  ])
```

In [39]:

```
# Accuracy, R^2s
print("Training accuracy: ",regressor.score(X_train,Y_train) )
print("Testing accuracy: ",regressor.score(X_test,Y_test) )
print("R^2 score: ",r2_score(Y_test,ypred) )
```

```
Training accuracy:  0.9527776419169073
Testing accuracy:  0.9150323289393305
r2 score:  0.9150323289393305
```

In [40]:

```
# Predict profit for { R&D Spend = 300000, Administration Spend = 160000, Marketing Spend = 200000 }
regressor.predict([[0,1,0,300000,160000,200000]])
```

Out[40]:

```
array([294953.71027927])
```

In [41]:

```
# Intercept Coefficient
print("Intercept: ", regressor.intercept_)
print("Coefficient: ", regressor.coef_)
```

```
Intercept:  47299.244832003154
Coefficient: [ 2.72986630e+02  2.14749984e+01 -2.94461629e+02  8.13783841
e-01
 -8.93750473e-03  2.46391951e-02]
```

In []: